

In [1]:

```

1 import pandas as pd
2 import numpy as np
3 import random
4
5 from sklearn.feature_extraction.text import CountVectorizer
6 from sklearn.feature_extraction.text import TfidfVectorizer
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.model_selection import train_test_split

```

In [2]:

```

1 from google.colab import files
2 uploaded = files.upload()

```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving urldata.csv to urldata.csv

In [3]:

```

1 urls_data = pd.read_csv("urldata.csv", index_col=0)
2 urls_data.head()

```

Out[3]:

	url	label	result
0	https://www.google.com	benign	0
1	https://www.youtube.com	benign	0
2	https://www.facebook.com	benign	0
3	https://www.baidu.com	benign	0
4	https://www.wikipedia.org	benign	0

In [4]:

```

1 def makeTokens(f):
2     tkns_BySlash = str(f.encode('utf-8')).split('/') # make tokens after splitting by slash
3     total_Tokens = []
4     for i in tkns_BySlash:
5         tokens = str(i).split('-') # make tokens after splitting by dash
6         tkns_ByDot = []
7         for j in range(0,len(tokens)):
8             temp_Tokens = str(tokens[j]).split('.') # make tokens after splitting by dot
9             tkns_ByDot = tkns_ByDot + temp_Tokens
10        total_Tokens = total_Tokens + tokens + tkns_ByDot
11    total_Tokens = list(set(total_Tokens)) #remove redundant tokens
12    if 'com' in total_Tokens:
13        total_Tokens.remove('com') #removing .com since it occurs a lot of times and it should not be included in our features
14    return total_Tokens

```

In [5]:

```
1 y = urls_data["result"]
```

In [6]:

```
1 y.unique()
```

Out[6]:

```
array([0, 1])
```

In [7]:

```
1 url_list = urls_data["url"]
```

In [8]:

```
1 vectorizer = TfidfVectorizer(tokenizer=makeTokens)
```

In [9]:

```
1 X = vectorizer.fit_transform(url_list)
```

In [10]:

```
1 X.shape
```

Out[10]:

(450176, 780471)

In [12]:

```
1 pip install catboost
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)

Collecting catboost

Downloading catboost-1.1.1-cp38-none-manylinux1_x86_64.whl (76.6 MB)

76.6/76.6 MB 11.2 MB/s eta 0:00:00

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.8/dist-packages (from catboost) (1.21.6)
 Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-packages (from catboost) (3.2.2)
 Requirement already satisfied: graphviz in /usr/local/lib/python3.8/dist-packages (from catboost) (0.10.1)
 Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (from catboost) (1.7.3)
 Requirement already satisfied: plotly in /usr/local/lib/python3.8/dist-packages (from catboost) (5.5.0)
 Requirement already satisfied: six in /usr/local/lib/python3.8/dist-packages (from catboost) (1.15.0)
 Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.8/dist-packages (from catboost) (1.3.5)
 Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-packages (from pandas>=0.24.0->catboost) (2022.7.1)
 Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.8/dist-packages (from pandas>=0.24.0->catboost) (2.8.2)
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (3.0.9)
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (1.4.4)
 Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (0.11.0)
 Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.8/dist-packages (from plotly->catboost) (8.1.0)
 Installing collected packages: catboost
 Successfully installed catboost-1.1.1

In [13]:

```
1 from sklearn.tree import DecisionTreeClassifier
2 from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier, BaggingClassifier
3 import xgboost as xgb
4 from catboost import CatBoostClassifier
5 from sklearn.svm import SVC
6 from sklearn.neighbors import KNeighborsClassifier
```

In [14]:

```
1 from sklearn.metrics import accuracy_score, confusion_matrix, recall_score, precision_score, f1_score
```

In [15]:

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In []:

```
1 logit = LogisticRegression()
2 logit.fit(X_train, y_train)
```

Out[18]:

LogisticRegression()

In []:

```
1 print("Accuracy ", logit.score(X_test, y_test))
```

Accuracy 0.9948020791683326

In []:

```
1 pred2 = logit.predict(X_test)
2 print(confusion_matrix(y_test, pred2))
3 print(accuracy_score(y_test, pred2), recall_score(y_test, pred2), precision_score(y_test, pred2), f1_score(y_test, pred2))
```

```
[[68919 2]
 [ 466 20649]]
0.9948020791683326 0.97793038124556 0.9999031523897148 0.9887947134032467
```

In []:

```

1 pred2 = logit.predict(X_train)
2 print(confusion_matrix(y_train, pred2))
3 print(accuracy_score(y_train, pred2), recall_score(y_train, pred2), precision_score(y_train, pred2), f1_score(y_train, pred2))

```

```

[[276813    4]
 [ 1409  81914]]
0.996076525795524 0.9830899031479904 0.999951170682878 0.9914488534927771

```

In []:

```

1 model2 = DecisionTreeClassifier()
2 model2.fit(X_train, y_train)
3
4 pred2 = model2.predict(X_test)
5 print(confusion_matrix(y_test, pred2))
6 print(accuracy_score(y_test, pred2), recall_score(y_test, pred2), precision_score(y_test, pred2), f1_score(y_test, pred2))

```

```

[[68864    57]
 [ 102 21013]]
0.9982340397174464 0.9951693109164101 0.9972947318462269 0.9962308877563114

```

In []:

```

1 pr2 = model2.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

```

[[276817    0]
 [    0  83323]]
1.0 1.0 1.0 1.0

```

In []:

```

1 model2 = DecisionTreeClassifier()
2 bag_clf = BaggingClassifier(base_estimator=model2, n_estimators=100)
3 bag_clf.fit(X_train, y_train)
4
5 pred_bag = bag_clf.predict(X_test)
6 print(confusion_matrix(y_test, pred_bag))
7 print(accuracy_score(y_test, pred_bag), recall_score(y_test, pred_bag), precision_score(y_test, pred_bag), f1_score(y_test, pred_bag))

```

In []:

```

1 pr2 = bag_clf.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

In [16]:

```
1 from sklearn.model_selection import GridSearchCV
```

In [21]:

```

1 gscv_rf = RandomForestClassifier(n_estimators=100, n_jobs=-1)
2 # param_rf = {'n_estimators':[200, 400, 500, 600, 700, 800, 1000]}
3
4 # gscv_rf = GridSearchCV(rf, param_grid=param_rf, n_jobs=-1)
5 gscv_rf.fit(X_train, y_train)
6
7 pred_rf = gscv_rf.predict(X_test)
8 print(confusion_matrix(y_test, pred_rf))
9 print(accuracy_score(y_test, pred_rf), recall_score(y_test, pred_rf), precision_score(y_test, pred_rf), f1_score(y_test, pred_rf))

```

```

[[68889    32]
 [ 119 20996]]
0.9983228930649962 0.9943641960691452 0.9984782195168347 0.9964169612984363

```

In [22]:

```

1 pr2 = gscv_rf.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

```

[[276817    0]
 [    0  83323]]
1.0 1.0 1.0 1.0

```

In [17]:

```

1 gscv_rf = RandomForestClassifier(n_estimators=50, n_jobs=-1)
2 # param_rf = {'n_estimators':[200, 400, 500, 600, 700, 800, 1000]}
3
4 # gscv_rf = GridSearchCV(rf, param_grid=param_rf, n_jobs=-1)
5 gscv_rf.fit(X_train, y_train)
6
7 pred_rf = gscv_rf.predict(X_test)
8 print(confusion_matrix(y_test, pred_rf))
9 print(accuracy_score(y_test, pred_rf), recall_score(y_test, pred_rf), precision_score(y_test, pred_rf), f1_score(y_test, pred_rf))

```

```

[[68887    34]
 [ 119 20996]]
0.9983006797281088 0.9943641960691452 0.9983832620066572 0.9963696761181635

```

In [18]:

```

1 pr2 = gscv_rf.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

```

[[276817    0]
 [    4 83319]]
0.9999888932081968 0.9999519940472619 1.0 0.9999759964474743

```

In []:

```
1 #gscv_rf.best_params_
```

In []:

```

1 gscv_gb = GradientBoostingClassifier(n_estimators=50, n_jobs=-1)
2 # param_gb = {'n_estimators':[700, 800, 900, 1000, 1100, 1200, 1300]}
3
4 # gscv_gb = GridSearchCV(gb, param_grid=param_gb, n_jobs=-1)
5 gscv_gb.fit(X_train, y_train)
6
7 pred_gb = gscv_gb.predict(X_test)
8 print(confusion_matrix(y_test, pred_gb))
9 print(accuracy_score(y_test, pred_gb), recall_score(y_test, pred_gb), precision_score(y_test, pred_gb), f1_score(y_test, pred_gb))

```

In []:

```

1 pr2 = gscv_gb.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

In []:

```
1 #gscv_gb.best_params_
```

In []:

```

1 gscv_ab = AdaBoostClassifier(n_estimators=50, n_jobs=-1)
2 # param_ab = {'n_estimators':[700, 800, 900, 1000, 1100, 1200, 1300]}
3
4 # gscv_ab = GridSearchCV(ab, param_grid=param_ab, n_jobs=-1)
5 gscv_ab.fit(X_train, y_train)
6
7 pred_ab = gscv_ab.predict(X_test)
8 print(confusion_matrix(y_test, pred_ab))
9 print(accuracy_score(y_test, pred_ab), recall_score(y_test, pred_ab), precision_score(y_test, pred_ab), f1_score(y_test, pred_ab))

```

In []:

```

1 pr2 = gscv_ab.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

In []:

```
1 #gscv_ab.best_params_
```

In []:

```

1 gscv_xb = xgb.XGBClassifier(n_estimators=500)
2 # param_xb = {'n_estimators':[700, 800, 900, 1000, 1100, 1200, 1300]}
3
4 # gscv_xb = GridSearchCV(xb, param_grid=param_xb, n_jobs=-1)
5 gscv_xb.fit(X_train, y_train)
6
7 pred_xb = gscv_xb.predict(X_test)
8 print(confusion_matrix(y_test, pred_xb))
9 print(accuracy_score(y_test, pred_xb), recall_score(y_test, pred_xb), precision_score(y_test, pred_xb), f1_score(y_test, pred_xb))

```

D:\AnacondaNavigator\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

warnings.warn(label_encoder_deprecation_msg, UserWarning)

[15:13:29] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```

[[68875   46]
 [ 134 20981]]
0.998000799680128 0.9936538006156761 0.9978123365197128 0.9957287266859666

```

In []:

```

1 pr2 = gscv_xb.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

```

[[276710   107]
 [   355 82968]]
0.9987171655467318 0.9957394716944901 0.9987120072223894 0.9972235243212058

```

In []:

```
1 #gscv_xb.best_params_
```

In []:

```

1 cb = CatBoostClassifier(n_estimators=500)
2 cb.fit(X_train, y_train)
3
4 pred_cb = cb.predict(X_test)
5 print(confusion_matrix(y_test, pred_cb))
6 print(accuracy_score(y_test, pred_cb), recall_score(y_test, pred_cb), precision_score(y_test, pred_cb), f1_score(y_test, pred_cb))

```

```

Learning rate set to 0.24021
0:   learn: 0.1891841      total: 4.92s   remaining: 40m 54s
1:   learn: 0.0704305      total: 8.63s   remaining: 35m 49s
2:   learn: 0.0357784      total: 11.8s   remaining: 32m 40s
3:   learn: 0.0234963      total: 15.3s   remaining: 31m 40s
4:   learn: 0.0185126      total: 18.8s   remaining: 30m 57s
5:   learn: 0.0163698      total: 22.1s   remaining: 30m 19s
6:   learn: 0.0158579      total: 25.3s   remaining: 29m 43s
7:   learn: 0.0155386      total: 28.5s   remaining: 29m 9s
8:   learn: 0.0151340      total: 31.6s   remaining: 28m 46s
9:   learn: 0.0149476      total: 34.9s   remaining: 28m 28s
10:  learn: 0.0148453      total: 38.2s   remaining: 28m 17s
11:  learn: 0.0139928      total: 41.6s   remaining: 28m 10s
12:  learn: 0.0139131      total: 44.7s   remaining: 27m 55s
13:  learn: 0.0138438      total: 47.9s   remaining: 27m 41s
14:  learn: 0.0137769      total: 51s     remaining: 27m 28s
15:  learn: 0.0137220      total: 54.1s   remaining: 27m 16s
16:  learn: 0.0136613      total: 57.2s   remaining: 27m 5s
17:  learn: 0.0136081      total: 1m      remaining: 26m 56s

```

In []:

```

1 pr2 = cb.predict(X_train)
2 print(confusion_matrix(y_train, pr2))
3 print(accuracy_score(y_train, pr2), recall_score(y_train, pr2), precision_score(y_train, pr2), f1_score(y_train, pr2))

```

```

[[276582   235]
 [   465 82858]]
0.9980563114344422 0.9944193079941913 0.9971718435969336 0.9957936736852225

```

In []:

```

1 import math
2 n = (int(math.sqrt(X_train.shape[0]))//2)*2+1

```

In []:

```

1 knn = KNeighborsClassifier(n_neighbors=n, weights='distance')
2 knn.fit(X_train, y_train)
3
4 pred_knn = knn.predict(X_test)
5 print(confusion_matrix(y_test, pred_knn))
6 print(accuracy_score(y_test, pred_knn), recall_score(y_test, pred_knn), precision_score(y_test, pred_knn), f1_score(y_test, pred_knn))

```

```

[[68078  843]
 [ 2477 18638]]
0.9631258607668044 0.882690030783803 0.9567270673990041 0.9182185436988866

```

In []:

```

1 sv = SVC()
2 sv.fit(X_train, y_train)
3
4 pred_sv = sv.predict(X_test)
5 print(confusion_matrix(y_test, pred_sv))
6 print(accuracy_score(y_test, pred_sv), recall_score(y_test, pred_sv), precision_score(y_test, pred_sv), f1_score(y_test, pred_sv))

```

In []:

```
1 from sklearn.naive_bayes import MultinomialNB, GaussianNB, BernoulliNB
```

In []:

```

1 mnb = MultinomialNB()
2 mnb.fit(X_train, y_train)
3
4 pred_nb = mnb.predict(X_test)
5 print(confusion_matrix(y_test, pred_nb))
6 print(accuracy_score(y_test, pred_nb), recall_score(y_test, pred_nb), precision_score(y_test, pred_nb), f1_score(y_test, pred_nb))

```

```

[[68860  61]
 [ 855 20260]]
0.98982629170554 0.9595074591522614 0.9969981792234635 0.9778936190752003

```

In []:

```

1 bnb = BernoulliNB()
2 bnb.fit(X_train, y_train)
3
4 pred_nb = bnb.predict(X_test)
5 print(confusion_matrix(y_test, pred_nb))
6 print(accuracy_score(y_test, pred_nb), recall_score(y_test, pred_nb), precision_score(y_test, pred_nb), f1_score(y_test, pred_nb))

```

```

[[68914  7]
 [ 327 20788]]
0.9962903727397929 0.9845133791143736 0.9996633806203414 0.9920305416368408

```

In []:

```

1 pred_nb = bnb.predict(X_train)
2 print(confusion_matrix(y_train, pred_nb))
3 print(accuracy_score(y_train, pred_nb), recall_score(y_train, pred_nb), precision_score(y_train, pred_nb), f1_score(y_train, pred_nb))

```

```

[[276805  12]
 [ 989 82334]]
0.9972205253512523 0.988130528185495 0.9998542734316178 0.993957831579837

```

In []:

```

1 lr_m = [99.48]
2 dt_m = [99.82]
3 xgb_m = [99.80]
4 knn_m = [96.31]
5 cb_m = [99.75]
6 mnb_m = [98.98]
7 bnb_m = [99.63]

```

In []:

```

1 for i in [53, 72, 96, 21]:
2     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=i, stratify=y)
3
4     model = LogisticRegression()
5     model.fit(X_train, y_train)
6
7     pred = model.predict(X_test)
8     print("LR:")
9     print(confusion_matrix(y_test, pred))
10    print(accuracy_score(y_test, pred), recall_score(y_test, pred), precision_score(y_test, pred), f1_score(y_test, pred))
11    lr_m.append(accuracy_score(y_test, pred))
12
13
14    model2 = DecisionTreeClassifier()
15    model2.fit(X_train, y_train)
16
17    pred2 = model2.predict(X_test)
18    print("DTC:")
19    print(confusion_matrix(y_test, pred2))
20    print(accuracy_score(y_test, pred2), recall_score(y_test, pred2), precision_score(y_test, pred2), f1_score(y_test, pred2))
21    dt_m.append(accuracy_score(y_test, pred2))
22
23
24    xb = xgb.XGBClassifier(n_estimators=700)
25    xb.fit(X_train, y_train)
26
27    pred_xb = xb.predict(X_test)
28    print("XGB:")
29    print(confusion_matrix(y_test, pred_xb))
30    print(accuracy_score(y_test, pred_xb), recall_score(y_test, pred_xb), precision_score(y_test, pred_xb), f1_score(y_test, pred_xb))
31    xgb_m.append(accuracy_score(y_test, pred_xb))
32
33
34    n = (int(math.sqrt(X_train.shape[0]))//2)*2+1
35
36    knn = KNeighborsClassifier(n_neighbors=n, weights='distance')
37    knn.fit(X_train, y_train)
38
39    pred_knn = knn.predict(X_test)
40    print("KNN:")
41    print(confusion_matrix(y_test, pred_knn))
42    print(accuracy_score(y_test, pred_knn), recall_score(y_test, pred_knn), precision_score(y_test, pred_knn), f1_score(y_test, pred_knn))
43    knn_m.append(accuracy_score(y_test, pred_knn))
44
45
46    cb = CatBoostClassifier(n_estimators=500)
47    cb.fit(X_train, y_train)
48
49    pred_cb = cb.predict(X_test)
50    print("CB:")
51    print(confusion_matrix(y_test, pred_cb))
52    print(accuracy_score(y_test, pred_cb), recall_score(y_test, pred_cb), precision_score(y_test, pred_cb), f1_score(y_test, pred_cb))
53    cb_m.append(accuracy_score(y_test, pred_cb))
54
55
56    mnb = MultinomialNB()
57    mnb.fit(X_train, y_train)
58
59    pred_nb = mnb.predict(X_test)
60    print("MNB:")
61    print(confusion_matrix(y_test, pred_nb))
62    print(accuracy_score(y_test, pred_nb), recall_score(y_test, pred_nb), precision_score(y_test, pred_nb), f1_score(y_test, pred_nb))
63    mnb_m.append(accuracy_score(y_test, pred_nb))
64
65
66    bnb = BernoulliNB()
67    bnb.fit(X_train, y_train)
68
69    pred_nb = bnb.predict(X_test)
70    print("BNB:")
71    print(confusion_matrix(y_test, pred_nb))
72    print(accuracy_score(y_test, pred_nb), recall_score(y_test, pred_nb), precision_score(y_test, pred_nb), f1_score(y_test, pred_nb))
73    bnb_m.append(accuracy_score(y_test, pred_nb))

```

D:\AnacondaNavigator\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

```
LR:
[[69147      1]
 [  435 20453]]
0.9951574925585321 0.9791746457296056 0.9999511098073727 0.9894538241981521
DTC:
[[69082      66]
 [  101 20787]]
```

In []:

```
1 for model in [lr_m, dt_m, xgb_m, cb_m, knn_m, mnb_m, bnb_m]:
2     print(np.mean(model), np.std(model))
```

99.478 0.019390719429666442
99.818 0.007483314773545635
99.792 0.007483314773545634
99.76599999999999 0.010198039027184655
96.28 0.060991802727905275
98.95 0.029664793948382496
99.636 0.012000000000002349

In []:

```
1
```